

Send all 3 fields to gemini-flash: the best reader on this sample, for \$4.29 instead of \$100 per 1,000 documents.

The 100 real receipts of the CORD v2 test split (Clova AI, CC BY 4.0), with CORD's own labels; the split has been public since 2022, so the vendors' models may have seen it. Both vendor prices were declared by Crusetra for this sample. File cord-labels-grouped.csv, 100 cases, SHA-256 e8ff6b84e865bef90dbc477a39727a78.... Measured 2026-09-29 22:49 UTC with tool commit 7e82950, 8 sources on 3 fields. Every figure below is read from the sealed record ac7d0adbe4907caf (a content hash).

\$95,710

saved a year at 1,000,000 pages, on your prices and volume

96.8 %

of total values read right by gemini-flash, against 93.7 % today

2 of 3

picks this sample cannot tell apart from another source within your margin

01

The routing, field by field

FIELD	SEND TO	RIGHT, 95 % INTERVAL, CASES	TODAY	AGAINST GOOGLE-EXPENSE, YOUR CURRENT CHAIN
total	gemini-flash	96.8 % [91-99], 95 cases	93.7 %	too close to call: 5 of 7 disagreements, p = 0.453
subtotal	gemini-flash	95.4 % [87-98], 65 cases	89.2 %	too close to call: 6 of 8 disagreements, p = 0.289
tax	gemini-flash	97.5 % [87-100], 40 cases	70 %	better: 12 of 13 disagreements, p = 0.003

02

Every source on every field: share of values read right, with its 95 % interval; the pick for each field is in green

SOURCE	COST	TOTAL	SUBTOTAL	TAX
gemini-flash	\$4.29 per 1,000 documents	96.8 [91-99]	95.4 [87-98]	97.5 [87-100]
google-expense	CURRENT \$100 per 1,000 documents	93.7 [87-97]	89.2 [79-95]	70 [55-82]
gen-4b	local, 804 ms a value	81.1 [72-88]	78.5 [67-87]	27.5 [16-43]
gen-0.6b	local, 315 ms a value	77.9 [69-85]	40 [29-52]	27.5 [16-43]
rules	local, under 1 ms a value	66.3 [56-75]	56.9 [45-68]	67.5 [52-80]
large	local, 21 ms a value	62.1 [52-71]	32.3 [22-44]	10 [4-23]
gen-8b	local, 1,162 ms a value	61.1 [51-70]	50.8 [39-63]	35 [22-50]
small	local, 14 ms a value	2.1 [1-7]	0 [0-6]	0 [0-9]

03

What the money rests on

At 1,000,000 pages a year (1 page a document), your current chain google-expense costs \$100,000; the routing above costs \$4,290. The saving, **\$95,710 a year**, rests on these inputs: google-expense \$100 per 1,000 documents, declared by Crusetra, for this sample; gemini-flash \$4.29 per 1,000 documents, declared by Crusetra, for this sample; volume 1,000,000 pages a year, declared. Local sources: machine time at \$1.20 an hour, an assumption; declare yours and the report follows. Accuracy and the verdicts are measured. The money is not: change a price and the figure moves with it.

04

What this sample cannot tell

These picks win on this sample, but the sample is too small to tell them apart; a larger sample could reverse them: gemini-flash against google-expense (your current chain) on total (95 cases, 7 disagreements, p = 0.453); subtotal (65 cases, 8 disagreements, p = 0.289).

05

How it was measured

Each source's answer is graded against the expected value, case by case, read as an amount written with grouped thousands (60.000 is sixty thousand). A case with no expected value is graded by nobody (total 5, subtotal 35, tax 60). Two sources are compared on the same cases: the cases where exactly one is right decide (McNemar, exact), and the worst case at 95 % is Newcombe's paired bound. A cheaper source is picked only if it cannot be more than 2 points worse than the best, the margin you declared. No rate is quoted under 20 cases. The local sources were asked a question made from each column name, such as "What is the total?"; your own wording can change their rates. The vendors' outputs are graded as they came.

06

What this report does not do

It does not check a price against an invoice: every price is the one declared above. It does not see a value: the record holds right, wrong or blank per case, never what was read. It measures these 100 cases; documents of another kind, layout or language need their own sample. Local sources were timed on the machine that ran the audit; another machine takes another time.